Loop Transformer

One shared layer stack. Dropout during training. Measurements at every loop.

A Reuse the same six layers

Weights stay shared; training masks are sampled again.

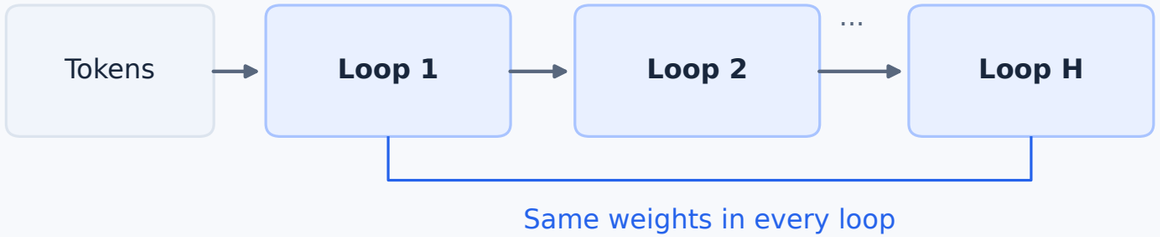

Inside each shared Transformer layer

Attention branch

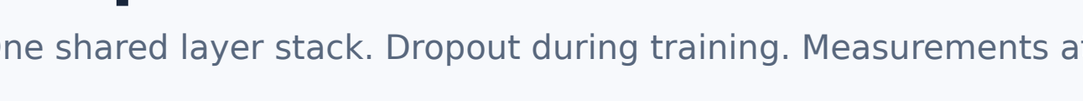

Then the FFN branch uses the same order.

Enable dropout on attention, FFN, or both.

TRAINING	EVALUATION
New dropout masks each call	Dropout off; full dense model

B Prediction quality

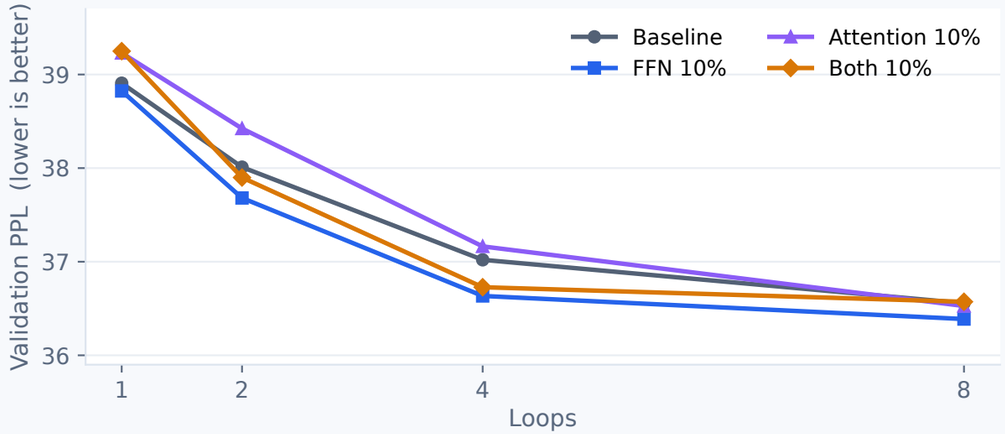

C FFN outputs, layer by layer

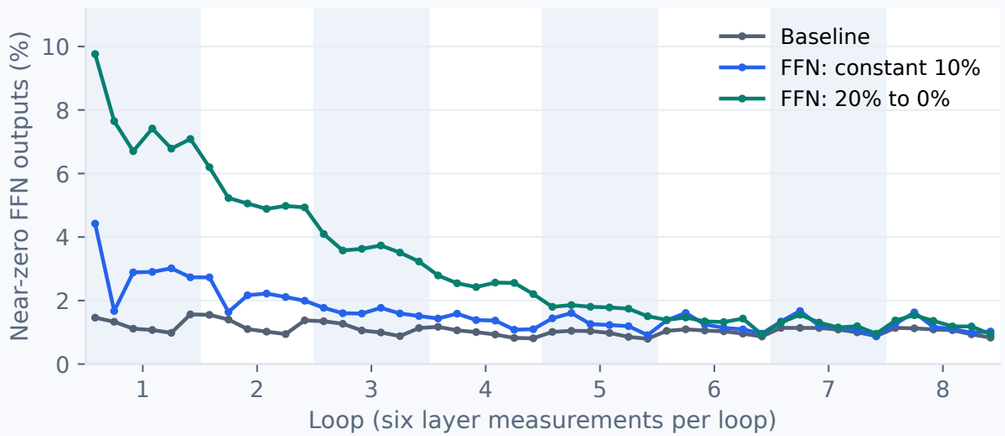

RECORDED RESULTS 80M parameters | 100,000 steps | 1.6384B training tokens | seed 42

Near-zero threshold: 0.01. Diagnostics use 512 fixed tokens with dropout off. Curves show measurements, not skipped computation.